

# B M S COLLEGE OF ENGINEERING

(An Autonomous Institution Affiliated to VTU, Belagavi)

Post Box No.: 1908, Bull Temple Road, Bengaluru – 560 019

## DEPARTMENT OF MACHINE LEARNING

Academic Year: 2023-2024 (Session: Nov 2023 - Mar 2024)



## INTRODUCTION TO NEURAL NETWORKS (23AM3PCINN)

ALTERNATIVE ASSESSMENT TOOL (AAT)

### MEDICAL IMAGE SEGMENTATION

#### Submitted by

Student Name:	Chetna Mundra	Aryaman Sharma
USN:	1BM21AI036	1BM21AI027
Date:	10-01-2024	
Semester & Section:	V-A	
Total Pages:	10	
Student Signature:		

#### Valuation Report (to be filled by the faculty)

Score:	
Comments:	
Faculty In-charge:	Dr. Seemanthini K
Faculty Signature: with date	

## INDEX

<b>CH. NO.</b>	<b>TITLE</b>	<b>PG. NO.</b>
1	Introduction	3
2	Methodology	4
	2.1 Dataset Preprocessing & Visualization	4
	2.2 Dealing with class imbalance	5
	2.3 Training	5
	2.4 Plotting validation accuracy and loss	5
	2.5 Estimation of classification performance	6
3	Flowchart	8
5	Results	9
6	Applications	10
7	Conclusion	11
Appendix 1	Code and Output	

# **1. INTRODUCTION**

In recent years, medical image segmentation has emerged as a pivotal tool in the field of diagnostic radiology, playing a crucial role in the accurate and timely identification of various diseases. Among these, pneumonia is a leading respiratory ailment with significant global health implications. Rapid and precise detection of pneumonia in chest X-rays is imperative for timely intervention and effective patient care. This report delves into the advancements and challenges associated with medical image segmentation techniques employed in the diagnosis of pneumonia through the analysis of X-rayed lungs.

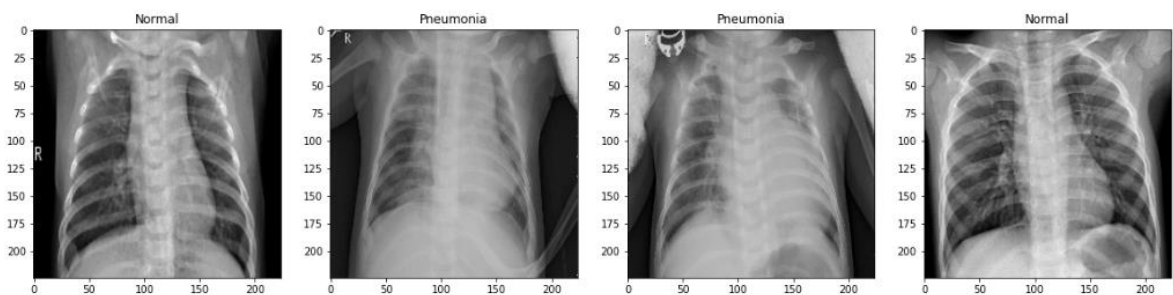
The conventional methods of pneumonia detection often rely on visual examination by radiologists, a process that is not only time-consuming but also subject to human error. The integration of artificial intelligence and image segmentation techniques has paved the way for more efficient and reliable diagnosis. By isolating and highlighting specific regions of interest within chest X-rays, these technologies facilitate the identification of pneumonia-related anomalies, thereby assisting healthcare professionals in making informed decisions and expediting patient treatment.

This report explores the current state-of-the-art in medical image segmentation for pneumonia diagnosis, encompassing the methodologies, challenges, and potential applications. Through an in-depth examination of the various techniques and technologies employed in this domain, we aim to shed light on the transformative impact of image segmentation on the accuracy and efficiency of pneumonia detection in chest X-rays. As we navigate through the intricacies of these cutting-edge approaches, a comprehensive understanding will be developed, laying the groundwork for improved diagnostic tools and enhanced patient outcomes.


## 2. METHODOLOGY

### 2.1 Dataset Preprocessing & Visualization

This section details the preparation and exploration of a chest X-ray dataset for pneumonia diagnosis. The dataset undergoes an 80%, 15%, 5% split for training, testing, and validation, respectively. Notably imbalanced, the dataset contains 1224 normal and 3418 pneumonia cases. Image loading and preprocessing involve resizing, converting to grayscale, transforming to RGB, and normalization. The resulting arrays, paired with labels (0 for normal, 1 for pneumonia), are created. Occurrence counts reveal the class imbalance in the training set. Visualizations provide insight into the preprocessed images, offering a representative glimpse of the dataset's diversity across training, testing, and validation sets. These steps lay the groundwork for subsequent model development, addressing imbalances and ensuring robust analysis in pneumonia detection.



*Fig 1. Refined Dataset*



*Fig 2. Graph Depicting Number of Images Alloted*

## 2.2 Dealing with class imbalance

To tackle the class imbalance, we adopt a pragmatic approach by assigning weights to each class, ensuring equitable learning for the Convolutional Neural Network (CNN). Utilizing sklearn's `compute_class_weight` function, class weights {0: 1.896, 1: 0.679} are computed, prioritizing the minority class (pneumonia). Additionally, to optimize memory usage, a balanced subset of images is created for training. The resulting datasets for training, testing, and validation, along with their respective shapes, are established. These measures strategically address the class imbalance, providing a foundation for a more effective CNN model training, where each class contributes proportionately to the learning process.

## 2.3 Training

The training phase commences with a batch size of 32, chosen for its efficiency and memory optimization. Prior to training, memory is cleared to enhance computational resources. The dataset lengths for training and validation are determined, and image augmentation techniques are employed to artificially expand the dataset, preventing overfitting.

A MobileNet architecture, a pre-trained Convolutional Neural Network (CNN), is utilized to expedite training. The model is compiled with binary crossentropy loss, Adam optimizer, and relevant evaluation metrics. To enhance the model's robustness, a class weight parameter is incorporated during training.

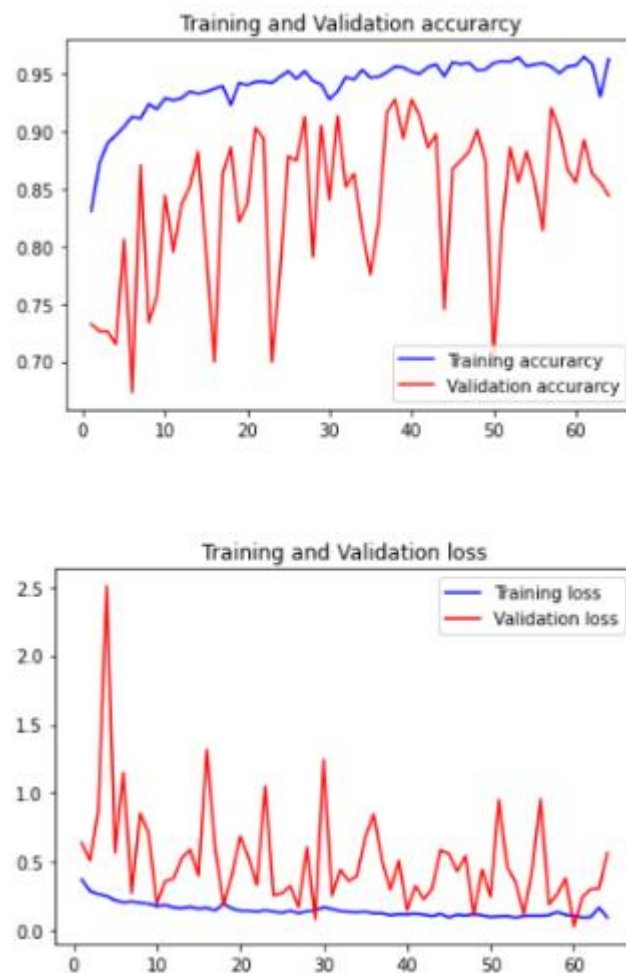
Training is conducted for 64 epochs using a generator-based approach, with image augmentation applied in real-time. The process leverages the power of transfer learning from MobileNet, showcasing its adaptability for pneumonia detection in chest X-rays. The training history is recorded for subsequent analysis, ensuring the model's efficacy in pneumonia classification.

## 2.4 Plotting validation accuracy and loss

The graph depicts the training and validation accuracy of a machine learning model over 60 epochs. The red line represents the training accuracy, which starts at 0.7 and increases to 0.93 by epoch 60. The blue line represents the validation accuracy, which starts at 0.7 and increases to 0.88 by epoch 60. The gap between the training and validation accuracy suggests that the model may be overfitting the training data.

In addition to the accuracy graph, the image also shows a graph of the training and validation loss. The training loss starts at 0.5 and decreases to 0.03 by epoch 60. The validation loss starts at 0.5 and decreases to 0.12 by epoch 60. The smaller gap between the training and validation loss compared to the accuracy suggests that the model may be underfitting the data.

Overall, the graphs suggest that the model is making progress on the training data, but it may not be generalizing well to unseen data. This could be due to overfitting or underfitting.




*Fig 3. Accuracy and Loss*

## 2.5 Estimation of classification performance

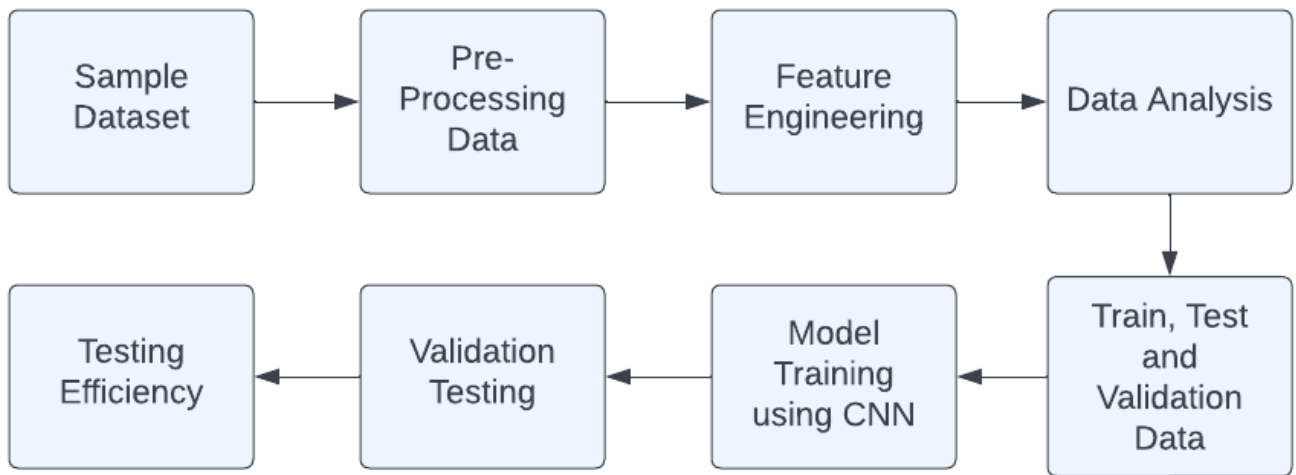
The model's classification performance is assessed using various metrics, including Accuracy, Precision, Recall, and F1-score. For the test set, the model achieves an Accuracy of 86.07%, Precision of 83.88%, Recall of 99.06%, and an F1-score of 90.84%. The confusion matrix reveals 156 true negatives, 122 false positives, 6 false negatives,

and 635 true positives. In the training phase, the model achieves an accuracy of 96.27%. The visual representation of the confusion matrix further illustrates the model's effectiveness in distinguishing between pneumonia and normal chest X-ray images.



*Fig 4. Heatmap*

### 3. FLOWCHART



*Fig 5. Flowchart*



## 4. RESULTS

ROC AUC Score: 0.96, indicating excellent model discrimination between pneumonia and normal chest X-ray images.

- Model Performance:
- Accuracy (test set): 86.07%
- Precision (test set): 83.88%
- Recall (test set): 99.06%
- F1-score (test set): 90.84%
- Accuracy (training set): 96.27%

AUC Score: 0.9572385773128768




Fig 6. ROC Curve

The ROC curve shows the trade-off between correctly identifying diseased patients (true positive rate) and incorrectly identifying healthy patients (false positive rate) at different threshold levels.

The AUC score of 0.96 indicates excellent performance of the model in distinguishing between pneumonia and normal chest X-ray images.

## 5. APPLICATIONS

The CNN model for pneumonia detection in chest X-rays has several valuable applications in the medical field:

1. **Early Diagnosis:** The model aids in the early detection of pneumonia, enabling prompt medical intervention and improving patient outcomes.
2. **Efficient Triage:** Emergency rooms and healthcare facilities can use the model to quickly triage patients with potential pneumonia, prioritizing those who require immediate attention.
3. **Resource Optimization:** Hospitals can optimize resources by swiftly identifying pneumonia cases, streamlining patient management, and allocating medical staff and equipment effectively.
4. **Telemedicine Support:** In remote or underserved areas, the model facilitates telemedicine by providing preliminary pneumonia assessments through chest X-ray scans, bridging the gap in healthcare access.
5. **Clinical Decision Support:** Healthcare professionals can use the model as a decision support tool, enhancing diagnostic accuracy and aiding in treatment planning.
6. **Research and Epidemiology:** The model contributes to large-scale studies and epidemiological research by efficiently analyzing chest X-rays, helping to understand the prevalence and trends of pneumonia.
7. **Educational Tool:** The model serves as an educational tool for medical students and practitioners, offering insights into image interpretation and enhancing diagnostic skills in pneumonia identification.

## **6. CONCLUSION**

The model demonstrates a strong ability to distinguish between pneumonia and normal chest X-ray images, achieving a high ROC AUC score of 0.96 and generally favorable performance metrics across accuracy, precision, recall, and F1-score. While exhibiting a slight decrease in accuracy from training to test sets, it maintains a high recall, suggesting a low rate of false negatives (misclassifying pneumonia cases as normal). This model shows significant potential for aiding in the accurate diagnosis of pneumonia.

```
In [1]: import os
import numpy as np
import pandas as pd
import pathlib
import imageio
```

```
In [2]: # Exploring dataset
base_dir = '../input/chest-xray-pneumonia/chest_xray/'

train_pneumonia_dir = base_dir+'train/PNEUMONIA/'
train_normal_dir=base_dir+'train/NORMAL/'

test_pneumonia_dir = base_dir+'test/PNEUMONIA/'
test_normal_dir = base_dir+'test/NORMAL/'

val_normal_dir= base_dir+'val/NORMAL/'
val_pnrmonia_dir= base_dir+'val/PNEUMONIA/'

train_pn = [train_pneumonia_dir+"{}".format(i) for i in os.listdir(train_pneumonia_dir)]
train_normal = [train_normal_dir+"{}".format(i) for i in os.listdir(train_normal_dir)]

test_normal = [test_normal_dir+"{}".format(i) for i in os.listdir(test_normal_dir)]
test_pn = [test_pneumonia_dir+"{}".format(i) for i in os.listdir(test_pneumonia_dir)]

val_pn= [val_pnrmonia_dir+"{}".format(i) for i in os.listdir(val_pnrmonia_dir) ]
val_normal= [val_normal_dir+"{}".format(i) for i in os.listdir(val_normal_dir) ]

print ("Total images:",len(train_pn+train_normal+test_normal+test_pn+val_pn+val_normal))
print ("Total pneumonia images:",len(train_pn+test_pn+val_pn))
print ("Total Nomral images:",len(train_normal+test_normal+val_normal))

Total images: 5856
Total pneumonia images: 4273
Total Nomral images: 1583
```

## Dataset Preprocessing & Visualization

```
In [3]: # Gathering all pneumina and normal chest X-ray in two python List
pn = train_pn + test_pn + val_pn
normal = train_normal + test_normal + val_normal

# Spliting dataset in train set,test set and validation set.

train_imgs = pn[:3418]+ normal[:1224] # 80% of 4273 Pneumonia and normal chest X-ray
test_imgs = pn[3418:4059]+ normal[1224:1502]
val_imgs = pn[4059:] + normal[1502:]

print("Total Train Images %s containing %s pneumonia and %s normal images"
      % (len(train_imgs),len(pn[:3418]),len(normal[:1224])))
print("Total Test Images %s containing %s pneumonia and %s normal images"
      % (len(test_imgs),len(pn[3418:4059]),len(normal[1224:1502])))
print("Total validation Images %s containing %s pneumonia and %s normal images"
      % (len(val_imgs),len(pn[4059:]),len(normal[1502:])))

import random

random.shuffle(train_imgs)
random.shuffle(test_imgs)
random.shuffle(val_imgs)
```

Total Train Images 4642 containing 3418 pneumonia and 1224 normal images  
Total Test Images 919 containing 641 pneumonia and 278 normal images  
Total validation Images 295 containing 214 pneumonia and 81 normal images

## Loading each image and their label into array

```
In [5]: import cv2
img_size = 224

def preprocess_image(image_list):

    X = []
    y = []
    count=0

    for image in image_list:

        try:

            img = cv2.imread(image,cv2.IMREAD_GRAYSCALE)

            img=cv2.resize(img,(img_size,img_size),interpolation=cv2.INTER_CUBIC)

            #convert image to 2D to 3D
            img = np.dstack([img, img, img])

            #convrt greyscale image to RGB
            img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

            # Normalalize Image
            img = img.astype(np.float32)/255.

            count=count+1

            X.append(img)

        except:
            continue
        #get the Labels
        if 'NORMAL' in image:
            y.append(0)

        elif 'IM' in image:
            y.append(0)

        elif 'virus' or 'bacteria' in image:
            y.append(1)

    return X, y
```

```
In [6]: X, y = preprocess_image(train_imgs)
```

```
In [7]: arr=y
uniqueValues, occurCount = np.unique(arr, return_counts=True)

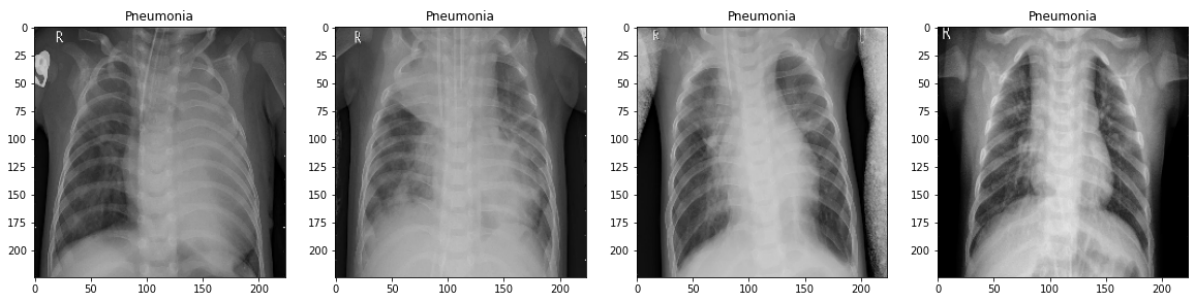
print("Unique Values : ", uniqueValues)
print("Occurrence Count : ", occurCount)
```

```
Unique Values : [0 1]
Occurrence Count : [1224 3418]
```

```
In [8]: import matplotlib.pyplot as plt

fig = plt.figure(figsize=(20, 5))
k=1
for i in range(4):
    a = fig.add_subplot(1, 4, k)
    if (y[i]==0):
        a.set_title('Normal')
    else:
        a.set_title('Pneumonia')

    plt.imshow(X[i])
    k=k+1;
```



```
In [9]: P, t = preprocess_image(test_imgs)
```

```
In [10]: arr=t
uniqueValues, occurCount = np.unique(arr, return_counts=True)

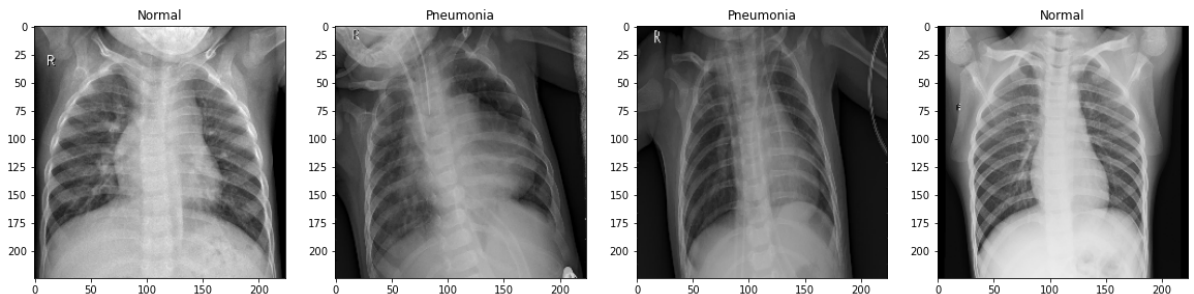
print("Unique Values : " , uniqueValues)
print("Occurrence Count : ", occurCount)

Unique Values : [0 1]
Occurrence Count : [278 641]
```

```
In [11]: import matplotlib.pyplot as plt

fig = plt.figure(figsize=(20, 5))
k=1
for i in range(4):
    a = fig.add_subplot(1, 4, k)
    if (t[i]==0):
        a.set_title('Normal')
    else:
        a.set_title('Pneumonia')

    plt.imshow(P[i])
    k=k+1;
```



```
In [12]: K, m = preprocess_image(val_imgs)
```

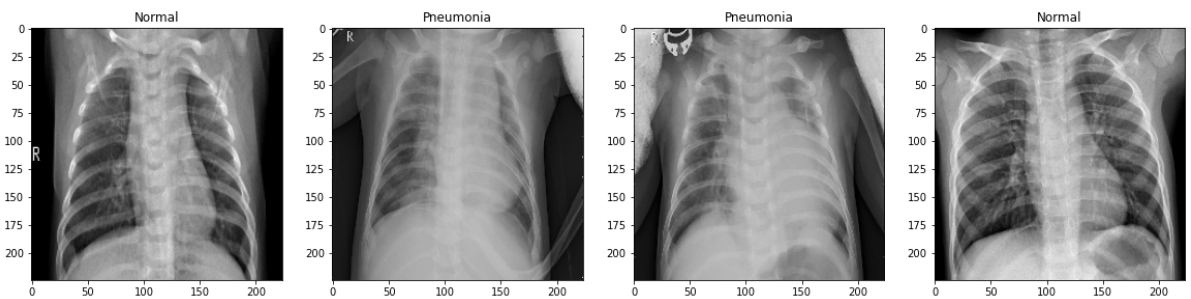
```
In [13]: arr=m
```

```
uniqueValues, occurCount = np.unique(arr, return_counts=True)
```


```
print("Unique Values : " , uniqueValues)  
print("Occurrence Count : ", occurCount)
```

```
Unique Values : [0 1]  
Occurrence Count : [ 81 214]
```

```
In [14]: import matplotlib.pyplot as plt  
  
fig = plt.figure(figsize=(20, 5))  
k=1  
for i in range(4):  
    a = fig.add_subplot(1, 4, k)  
    if (m[i]==0):  
        a.set_title('Normal')  
    else:  
        a.set_title('Pneumonia')  
  
    plt.imshow(K[i])  
    k=k+1;
```



```
In [15]: import seaborn as sns  
  
df=pd.DataFrame()  
df['Train']=y  
df['Test']=pd.Series(t)  
df['Val']=pd.Series(m)  
  
fig, ax =plt.subplots(1,3,figsize=(15,7))  
sns.countplot(df['Train'], ax=ax[0])  
ax[0].set(ylim=(0, 3500))  
  
sns.countplot(df['Test'], ax=ax[1])  
ax[1].set(ylim=(0, 3500))  
  
sns.countplot(df['Val'], ax=ax[2])  
ax[2].set(ylim=(0, 3500))  
  
fig.show()
```



```
In [16]: from sklearn.utils import class_weight

class_weights = class_weight.compute_class_weight('balanced',
                                                    np.unique(y), # here, y contains the labels
                                                    y)

class_weights = dict(enumerate(class_weights))
print(class_weights)

{0: 1.8962418300653594, 1: 0.679052077238151}
```

```
In [17]: import seaborn as sns
import gc

train_imgs = train_pn[:3875] + train_normal[:1341]
del train_imgs
gc.collect()

X_train = np.array(X)
y_train = np.array(y)
X_test = np.array(P)
y_test = np.array(t)
X_val = np.array(K)
y_val = np.array(m)

print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
print(X_val.shape)
print(y_val.shape)

(4642, 224, 224, 3)
(4642,)
(919, 224, 224, 3)
(919,)
(295, 224, 224, 3)
(295,)
```

## Training

```
In [18]: # clear memory
del X
del y
gc.collect()
```



```

#get the length of the train and validation data
ntrain = len(X_train)
nval = len(X_val)

batch_size = 32

```

```

In [19]: from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator( rotation_range=7,
                                     width_shift_range=0.05,
                                     height_shift_range=0.05,
                                     shear_range=0.2,
                                     zoom_range=0.45,
                                     horizontal_flip=True)

val_datagen = ImageDataGenerator(zoom_range=0.45)

```

Using TensorFlow backend.

```

In [20]: train_generator = train_datagen.flow(X_train, y_train, batch_size=batch_size)
val_generator = val_datagen.flow(X_val, y_val, batch_size=batch_size)

```

```

In [21]: img_size = 224

```

```

In [22]: from keras import layers
from keras import models
from keras import optimizers
from keras.applications import *
from keras.layers import Dense, GlobalAveragePooling2D
from keras.preprocessing.image import img_to_array, load_img
from keras.models import Model
from keras import backend as K

base_model = MobileNet(weights=None, include_top=False, input_shape=(img_size, img_size, 3))

x = base_model.output

x = GlobalAveragePooling2D()(x)

predictions = Dense(1, activation="sigmoid")(x)

model = Model(inputs=base_model.input, outputs=predictions)
# Compile model
model.compile(optimizer='adam', loss = 'binary_crossentropy',
              metrics = ['binary_accuracy', 'mae'])

```

```

In [24]: history = model.fit_generator(train_generator,
                                       steps_per_epoch=ntrain // batch_size,
                                       epochs=64,
                                       validation_data=val_generator,
                                       validation_steps=nval // batch_size,
                                       class_weight =class_weights,
                                       )

```

Epoch 1/64  
145/145 [=====] - 82s 563ms/step - loss: 0.3733 - binary\_accuracy: 0.8308 - mae: 0.2146 - val\_loss: 0.6394 - val\_binary\_accuracy: 0.7326 - val\_mae: 0.3425

Epoch 2/64  
145/145 [=====] - 70s 485ms/step - loss: 0.2884 - binary\_accuracy: 0.8722 - mae: 0.1633 - val\_loss: 0.5110 - val\_binary\_accuracy: 0.7262 - val\_mae: 0.3222

Epoch 3/64  
145/145 [=====] - 70s 481ms/step - loss: 0.2651 - binary\_accuracy: 0.8898 - mae: 0.1498 - val\_loss: 0.8732 - val\_binary\_accuracy: 0.7262 - val\_mae: 0.2890

Epoch 4/64  
145/145 [=====] - 70s 481ms/step - loss: 0.2506 - binary\_accuracy: 0.8965 - mae: 0.1382 - val\_loss: 2.5055 - val\_binary\_accuracy: 0.7148 - val\_mae: 0.2860

Epoch 5/64  
145/145 [=====] - 70s 484ms/step - loss: 0.2218 - binary\_accuracy: 0.9039 - mae: 0.1311 - val\_loss: 0.5667 - val\_binary\_accuracy: 0.8061 - val\_mae: 0.2099

Epoch 6/64  
145/145 [=====] - 69s 475ms/step - loss: 0.2055 - binary\_accuracy: 0.9128 - mae: 0.1179 - val\_loss: 1.1499 - val\_binary\_accuracy: 0.6730 - val\_mae: 0.3128

Epoch 7/64  
145/145 [=====] - 69s 474ms/step - loss: 0.2121 - binary\_accuracy: 0.9108 - mae: 0.1212 - val\_loss: 0.2709 - val\_binary\_accuracy: 0.8707 - val\_mae: 0.1499

Epoch 8/64  
145/145 [=====] - 70s 481ms/step - loss: 0.2036 - binary\_accuracy: 0.9236 - mae: 0.1151 - val\_loss: 0.8545 - val\_binary\_accuracy: 0.7338 - val\_mae: 0.2577

Epoch 9/64  
145/145 [=====] - 69s 478ms/step - loss: 0.1978 - binary\_accuracy: 0.9191 - mae: 0.1107 - val\_loss: 0.7145 - val\_binary\_accuracy: 0.7567 - val\_mae: 0.2580

Epoch 10/64  
145/145 [=====] - 69s 479ms/step - loss: 0.1802 - binary\_accuracy: 0.9286 - mae: 0.1018 - val\_loss: 0.1946 - val\_binary\_accuracy: 0.8441 - val\_mae: 0.1726

Epoch 11/64  
145/145 [=====] - 69s 474ms/step - loss: 0.1884 - binary\_accuracy: 0.9269 - mae: 0.1037 - val\_loss: 0.3602 - val\_binary\_accuracy: 0.7951 - val\_mae: 0.2373

Epoch 12/64  
145/145 [=====] - 72s 494ms/step - loss: 0.1685 - binary\_accuracy: 0.9286 - mae: 0.0983 - val\_loss: 0.3782 - val\_binary\_accuracy: 0.8365 - val\_mae: 0.1905

Epoch 13/64  
145/145 [=====] - 68s 472ms/step - loss: 0.1632 - binary\_accuracy: 0.9345 - mae: 0.0922 - val\_loss: 0.5275 - val\_binary\_accuracy: 0.8517 - val\_mae: 0.1582

Epoch 14/64  
145/145 [=====] - 70s 482ms/step - loss: 0.1727 - binary\_accuracy: 0.9328 - mae: 0.0956 - val\_loss: 0.5903 - val\_binary\_accuracy: 0.8821 - val\_mae: 0.1404

Epoch 15/64  
145/145 [=====] - 70s 486ms/step - loss: 0.1587 - binary\_accuracy: 0.9343 - mae: 0.0919 - val\_loss: 0.4016 - val\_binary\_accuracy: 0.7947 - val\_mae: 0.2206

Epoch 16/64  
145/145 [=====] - 71s 489ms/step - loss: 0.1652 - binary\_accuracy: 0.9369 - mae: 0.0946 - val\_loss: 1.3163 - val\_binary\_accuracy: 0.6996 - val\_mae: 0.3074

Epoch 17/64  
145/145 [=====] - 70s 480ms/step - loss: 0.1470 - binary\_accuracy: 0.9393 - mae: 0.0848 - val\_loss: 0.5998 - val\_binary\_accuracy: 0.8631 - val\_mae: 0.1444

Epoch 18/64  
145/145 [=====] - 68s 471ms/step - loss: 0.2005 - binary\_accuracy: 0.9228 - mae: 0.1095 - val\_loss: 0.1908 - val\_binary\_accuracy: 0.8859 - val\_mae: 0.1341

Epoch 19/64  
145/145 [=====] - 69s 474ms/step - loss: 0.1618 - binary\_accuracy: 0.9418 - mae: 0.0873 - val\_loss: 0.4009 - val\_binary\_accuracy: 0.8213 - val\_mae: 0.2007

Epoch 20/64  
145/145 [=====] - 67s 462ms/step - loss: 0.1474 - binary\_accuracy: 0.9402 - mae: 0.0838 - val\_loss: 0.6864 - val\_binary\_accuracy: 0.8365 - val\_mae: 0.2042

Epoch 21/64  
145/145 [=====] - 69s 478ms/step - loss: 0.1453 - binary\_accuracy: 0.9429 - mae: 0.0820 - val\_loss: 0.5333 - val\_binary\_accuracy: 0.9028 - val\_mae: 0.1253

Epoch 22/64  
145/145 [=====] - 70s 481ms/step - loss: 0.1367 - binary\_accuracy: 0.9432 - mae: 0.0759 - val\_loss: 0.3326 - val\_binary\_accuracy: 0.8935 - val\_mae: 0.1294

Epoch 23/64  
145/145 [=====] - 68s 468ms/step - loss: 0.1486 - binary\_accuracy: 0.9419 - mae: 0.0828 - val\_loss: 1.0472 - val\_binary\_accuracy: 0.6996 - val\_mae: 0.3262

Epoch 24/64  
145/145 [=====] - 66s 458ms/step - loss: 0.1381 - binary\_accuracy: 0.9473 - mae: 0.0773 - val\_loss: 0.2562 - val\_binary\_accuracy: 0.7795 - val\_mae: 0.2285

Epoch 25/64  
145/145 [=====] - 67s 464ms/step - loss: 0.1289 - binary\_accuracy: 0.9521 - mae: 0.0717 - val\_loss: 0.2720 - val\_binary\_accuracy: 0.8783 - val\_mae: 0.1314

Epoch 26/64  
145/145 [=====] - 68s 469ms/step - loss: 0.1469 - binary\_accuracy: 0.9456 - mae: 0.0790 - val\_loss: 0.3243 - val\_binary\_accuracy: 0.8745 - val\_mae: 0.1518

Epoch 27/64  
145/145 [=====] - 68s 472ms/step - loss: 0.1263 - binary\_accuracy: 0.9522 - mae: 0.0714 - val\_loss: 0.1715 - val\_binary\_accuracy: 0.9125 - val\_mae: 0.1201

Epoch 28/64  
145/145 [=====] - 69s 478ms/step - loss: 0.1447 - binary\_accuracy: 0.9437 - mae: 0.0775 - val\_loss: 0.6071 - val\_binary\_accuracy: 0.7909 - val\_mae: 0.2516

Epoch 29/64  
145/145 [=====] - 69s 474ms/step - loss: 0.1481 - binary\_accuracy: 0.9408 - mae: 0.0822 - val\_loss: 0.0827 - val\_binary\_accuracy: 0.9049 - val\_mae: 0.1166

Epoch 30/64  
145/145 [=====] - 68s 468ms/step - loss: 0.2016 - binary\_accuracy: 0.9278 - mae: 0.1022 - val\_loss: 1.2444 - val\_binary\_accuracy: 0.8403 - val\_mae: 0.1715

Epoch 31/64  
145/145 [=====] - 69s 472ms/step - loss: 0.1574 - binary\_accuracy: 0.9347 - mae: 0.0912 - val\_loss: 0.2534 - val\_binary\_accuracy: 0.9132 - val\_mae: 0.1156

Epoch 32/64  
145/145 [=====] - 72s 495ms/step - loss: 0.1464 - binary\_accuracy: 0.9471 - mae: 0.0801 - val\_loss: 0.4451 - val\_binary\_accuracy: 0.8517 - val\_mae: 0.1649

Epoch 33/64  
145/145 [=====] - 68s 471ms/step - loss: 0.1351 - binary\_accuracy: 0.9447 - mae: 0.0759 - val\_loss: 0.3611 - val\_binary\_accuracy: 0.8631 - val\_mae: 0.1611

Epoch 34/64  
145/145 [=====] - 69s 473ms/step - loss: 0.1332 - binary\_accuracy: 0.9534 - mae: 0.0740 - val\_loss: 0.3995 - val\_binary\_accuracy: 0.8137 - val\_mae: 0.2037

Epoch 35/64  
145/145 [=====] - 70s 481ms/step - loss: 0.1371 - binary\_accuracy: 0.9464 - mae: 0.0782 - val\_loss: 0.6809 - val\_binary\_accuracy: 0.7757 - val\_mae: 0.2259

Epoch 36/64  
145/145 [=====] - 69s 477ms/step - loss: 0.1270 - binary\_accuracy: 0.9474 - mae: 0.0717 - val\_loss: 0.8468 - val\_binary\_accuracy: 0.8213 - val\_mae: 0.1749

Epoch 37/64  
145/145 [=====] - 68s 472ms/step - loss: 0.1257 - binary\_accuracy: 0.9512 - mae: 0.0690 - val\_loss: 0.5111 - val\_binary\_accuracy: 0.9163 - val\_mae: 0.1042

Epoch 38/64  
145/145 [=====] - 67s 464ms/step - loss: 0.1124 - binary\_accuracy: 0.9562 - mae: 0.0637 - val\_loss: 0.2978 - val\_binary\_accuracy: 0.9278 - val\_mae: 0.1044

Epoch 39/64  
145/145 [=====] - 67s 464ms/step - loss: 0.1197 - binary\_accuracy: 0.9555 - mae: 0.0666 - val\_loss: 0.5098 - val\_binary\_accuracy: 0.8935 - val\_mae: 0.1225

Epoch 40/64  
145/145 [=====] - 69s 476ms/step - loss: 0.1185 - binary\_accuracy: 0.9516 - mae: 0.0672 - val\_loss: 0.1492 - val\_binary\_accuracy: 0.9278 - val\_mae: 0.1012

Epoch 41/64  
145/145 [=====] - 69s 476ms/step - loss: 0.1237 - binary\_accuracy: 0.9499 - mae: 0.0685 - val\_loss: 0.3255 - val\_binary\_accuracy: 0.9132 - val\_mae: 0.1150

Epoch 42/64  
145/145 [=====] - 71s 492ms/step - loss: 0.1169 - binary\_accuracy: 0.9556 - mae: 0.0648 - val\_loss: 0.2300 - val\_binary\_accuracy: 0.8859 - val\_mae: 0.1278

Epoch 43/64  
145/145 [=====] - 70s 486ms/step - loss: 0.1055 - binary\_accuracy: 0.9581 - mae: 0.0596 - val\_loss: 0.3041 - val\_binary\_accuracy: 0.8973 - val\_mae: 0.1139

Epoch 44/64  
145/145 [=====] - 67s 464ms/step - loss: 0.1230 - binary\_accuracy: 0.9475 - mae: 0.0685 - val\_loss: 0.5850 - val\_binary\_accuracy: 0.7452 - val\_mae: 0.2874

Epoch 45/64  
145/145 [=====] - 68s 469ms/step - loss: 0.1049 - binary\_accuracy: 0.9601 - mae: 0.0563 - val\_loss: 0.5603 - val\_binary\_accuracy: 0.8669 - val\_mae: 0.1388

Epoch 46/64  
145/145 [=====] - 68s 467ms/step - loss: 0.1162 - binary\_accuracy: 0.9580 - mae: 0.0606 - val\_loss: 0.4321 - val\_binary\_accuracy: 0.8745 - val\_mae: 0.1485

Epoch 47/64  
145/145 [=====] - 70s 480ms/step - loss: 0.1100 - binary\_accuracy: 0.9596 - mae: 0.0622 - val\_loss: 0.5426 - val\_binary\_accuracy: 0.8821 - val\_mae: 0.1264

Epoch 48/64  
145/145 [=====] - 70s 480ms/step - loss: 0.1250 - binary\_accuracy: 0.9527 - mae: 0.0654 - val\_loss: 0.1165 - val\_binary\_accuracy: 0.9011 - val\_mae: 0.1217

Epoch 49/64  
145/145 [=====] - 68s 467ms/step - loss: 0.1117 - binary\_accuracy: 0.9537 - mae: 0.0640 - val\_loss: 0.4459 - val\_binary\_accuracy: 0.8745 - val\_mae: 0.1709  
Epoch 50/64  
145/145 [=====] - 66s 457ms/step - loss: 0.0992 - binary\_accuracy: 0.9592 - mae: 0.0567 - val\_loss: 0.2475 - val\_binary\_accuracy: 0.7034 - val\_mae: 0.3188  
Epoch 51/64  
145/145 [=====] - 66s 457ms/step - loss: 0.1039 - binary\_accuracy: 0.9610 - mae: 0.0563 - val\_loss: 0.9502 - val\_binary\_accuracy: 0.8160 - val\_mae: 0.1892  
Epoch 52/64  
145/145 [=====] - 68s 470ms/step - loss: 0.1064 - binary\_accuracy: 0.9603 - mae: 0.0589 - val\_loss: 0.4620 - val\_binary\_accuracy: 0.8859 - val\_mae: 0.1388  
Epoch 53/64  
145/145 [=====] - 66s 455ms/step - loss: 0.0956 - binary\_accuracy: 0.9642 - mae: 0.0535 - val\_loss: 0.3536 - val\_binary\_accuracy: 0.8555 - val\_mae: 0.1537  
Epoch 54/64  
145/145 [=====] - 66s 452ms/step - loss: 0.1080 - binary\_accuracy: 0.9564 - mae: 0.0607 - val\_loss: 0.1186 - val\_binary\_accuracy: 0.8821 - val\_mae: 0.1505  
Epoch 55/64  
145/145 [=====] - 65s 451ms/step - loss: 0.1089 - binary\_accuracy: 0.9579 - mae: 0.0603 - val\_loss: 0.4492 - val\_binary\_accuracy: 0.8555 - val\_mae: 0.1603  
Epoch 56/64  
145/145 [=====] - 65s 447ms/step - loss: 0.1086 - binary\_accuracy: 0.9590 - mae: 0.0599 - val\_loss: 0.9572 - val\_binary\_accuracy: 0.8137 - val\_mae: 0.1917  
Epoch 57/64  
145/145 [=====] - 64s 444ms/step - loss: 0.1168 - binary\_accuracy: 0.9562 - mae: 0.0637 - val\_loss: 0.1908 - val\_binary\_accuracy: 0.9202 - val\_mae: 0.1060  
Epoch 58/64  
145/145 [=====] - 65s 448ms/step - loss: 0.1481 - binary\_accuracy: 0.9503 - mae: 0.0767 - val\_loss: 0.2648 - val\_binary\_accuracy: 0.9011 - val\_mae: 0.1337  
Epoch 59/64  
145/145 [=====] - 63s 437ms/step - loss: 0.1160 - binary\_accuracy: 0.9563 - mae: 0.0644 - val\_loss: 0.3824 - val\_binary\_accuracy: 0.8669 - val\_mae: 0.1462  
Epoch 60/64  
145/145 [=====] - 65s 448ms/step - loss: 0.1093 - binary\_accuracy: 0.9569 - mae: 0.0614 - val\_loss: 0.0294 - val\_binary\_accuracy: 0.8555 - val\_mae: 0.1441  
Epoch 61/64  
145/145 [=====] - 65s 448ms/step - loss: 0.0968 - binary\_accuracy: 0.9646 - mae: 0.0523 - val\_loss: 0.2439 - val\_binary\_accuracy: 0.8924 - val\_mae: 0.1300  
Epoch 62/64  
145/145 [=====] - 67s 459ms/step - loss: 0.0979 - binary\_accuracy: 0.9581 - mae: 0.0544 - val\_loss: 0.3027 - val\_binary\_accuracy: 0.8631 - val\_mae: 0.1454  
Epoch 63/64  
145/145 [=====] - 64s 441ms/step - loss: 0.1866 - binary\_accuracy: 0.9301 - mae: 0.0979 - val\_loss: 0.3057 - val\_binary\_accuracy: 0.8555 - val\_mae: 0.1643  
Epoch 64/64  
145/145 [=====] - 64s 443ms/step - loss: 0.0954 - binary\_accuracy: 0.9627 - mae: 0.0532 - val\_loss: 0.5651 - val\_binary\_accuracy: 0.8441 - val\_mae: 0.1734

## Plot how validation accuracy and loss are increasing against training accuracy and loss.



```
In [25]: # Lets plot the train and val curve
# Get the details form the history object
acc = history.history['binary_accuracy']
val_acc = history.history['val_binary_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

#Train and validation accuracy
plt.plot(epochs, acc, 'b', label='Training accuracy')
plt.plot(epochs, val_acc, 'r', label='Validation accuracy')
plt.title('Training and Validation accuracy')
plt.legend()

plt.figure()
#Train and validation loss
plt.plot(epochs, loss, 'b', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and Validation loss')
plt.legend()

plt.show()
```



```
In [26]: from sklearn.metrics import accuracy_score, confusion_matrix

preds = model.predict(X_test)

acc = accuracy_score(y_test, np.round(preds))*100
cm = confusion_matrix(y_test, np.round(preds))

tn, fp, fn, tp = cm.ravel()

print('CONFUSION MATRIX -----')
print(cm)

print('\n=====TEST METRICS=====')
precision = tp/(tp+fp)*100
recall = tp/(tp+fn)*100
print('Accuracy: {}'.format(acc))
print('Precision: {}'.format(precision))
print('Recall: {}'.format(recall))
print('F1-score: {}'.format(2*precision*recall/(precision+recall)))

print('\nTRAIN METRIC -----')
print('Train acc: {}'.format(np.round((history.history['binary_accuracy'][-1])*100,


CONFUSION MATRIX -----
[[156 122]
 [  6 635]]

=====TEST METRICS=====
Accuracy: 86.07181719260065%
Precision: 83.88375165125495%
Recall: 99.06396255850234%
F1-score: 90.84406294706724

TRAIN METRIC -----
Train acc: 96.27
```

```
In [27]: import seaborn as sns
sns.heatmap(cm, annot=True, fmt="d",)
```

```
Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4b74f8c290>
```



```
In [28]: from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.metrics import auc

fpr , tpr , thresholds = roc_curve ( y_test , preds)
auc_keras = auc(fpr, tpr)
print("AUC Score:",auc_keras)
plt.figure()
```

```

lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='ROC curve (area = %0.2f)' % auc_keras)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

```

AUC Score: 0.9572385773128768

